

# Map Reduce

**: Simplified Data Processing on Large Clusters**

Jeffrey Dean and Sanjay Ghemawat


Google, Inc

presented by Joon-Hyun Jeong

# Outline

1. Intro
2. Implementation of Map Reduce
3. Refinements
4. Results
5. Conclusion

# Intro



# Motivation

- \* Processing large amounts of raw data in reasonable time**
  - How to parallelize the computation?
  - How to distribute the data?
  - How to handle machine failures?
- \* Observation: most google computations applied map, reduce operation**

# Where to go?

**Functional model with *user specified* map and reduce operations**

- easily parallelize large computations
- fault tolerance by re-execution mechanism

# Programming model with example

```
map(String key, String value):  
    // key: document name  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");
```

input to Map function: (k1, v1)  
output from Map function: list (k2, v2)

```
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```


Input to Reduce function: (k2, list (v2))  
Output from Reduce function: list(v2)

# Implementation of Map Reduce



# Implementation environment

**A Cluster has hundreds or thousands of commodity machines.**



**CPU: dual-processor x86 processors**

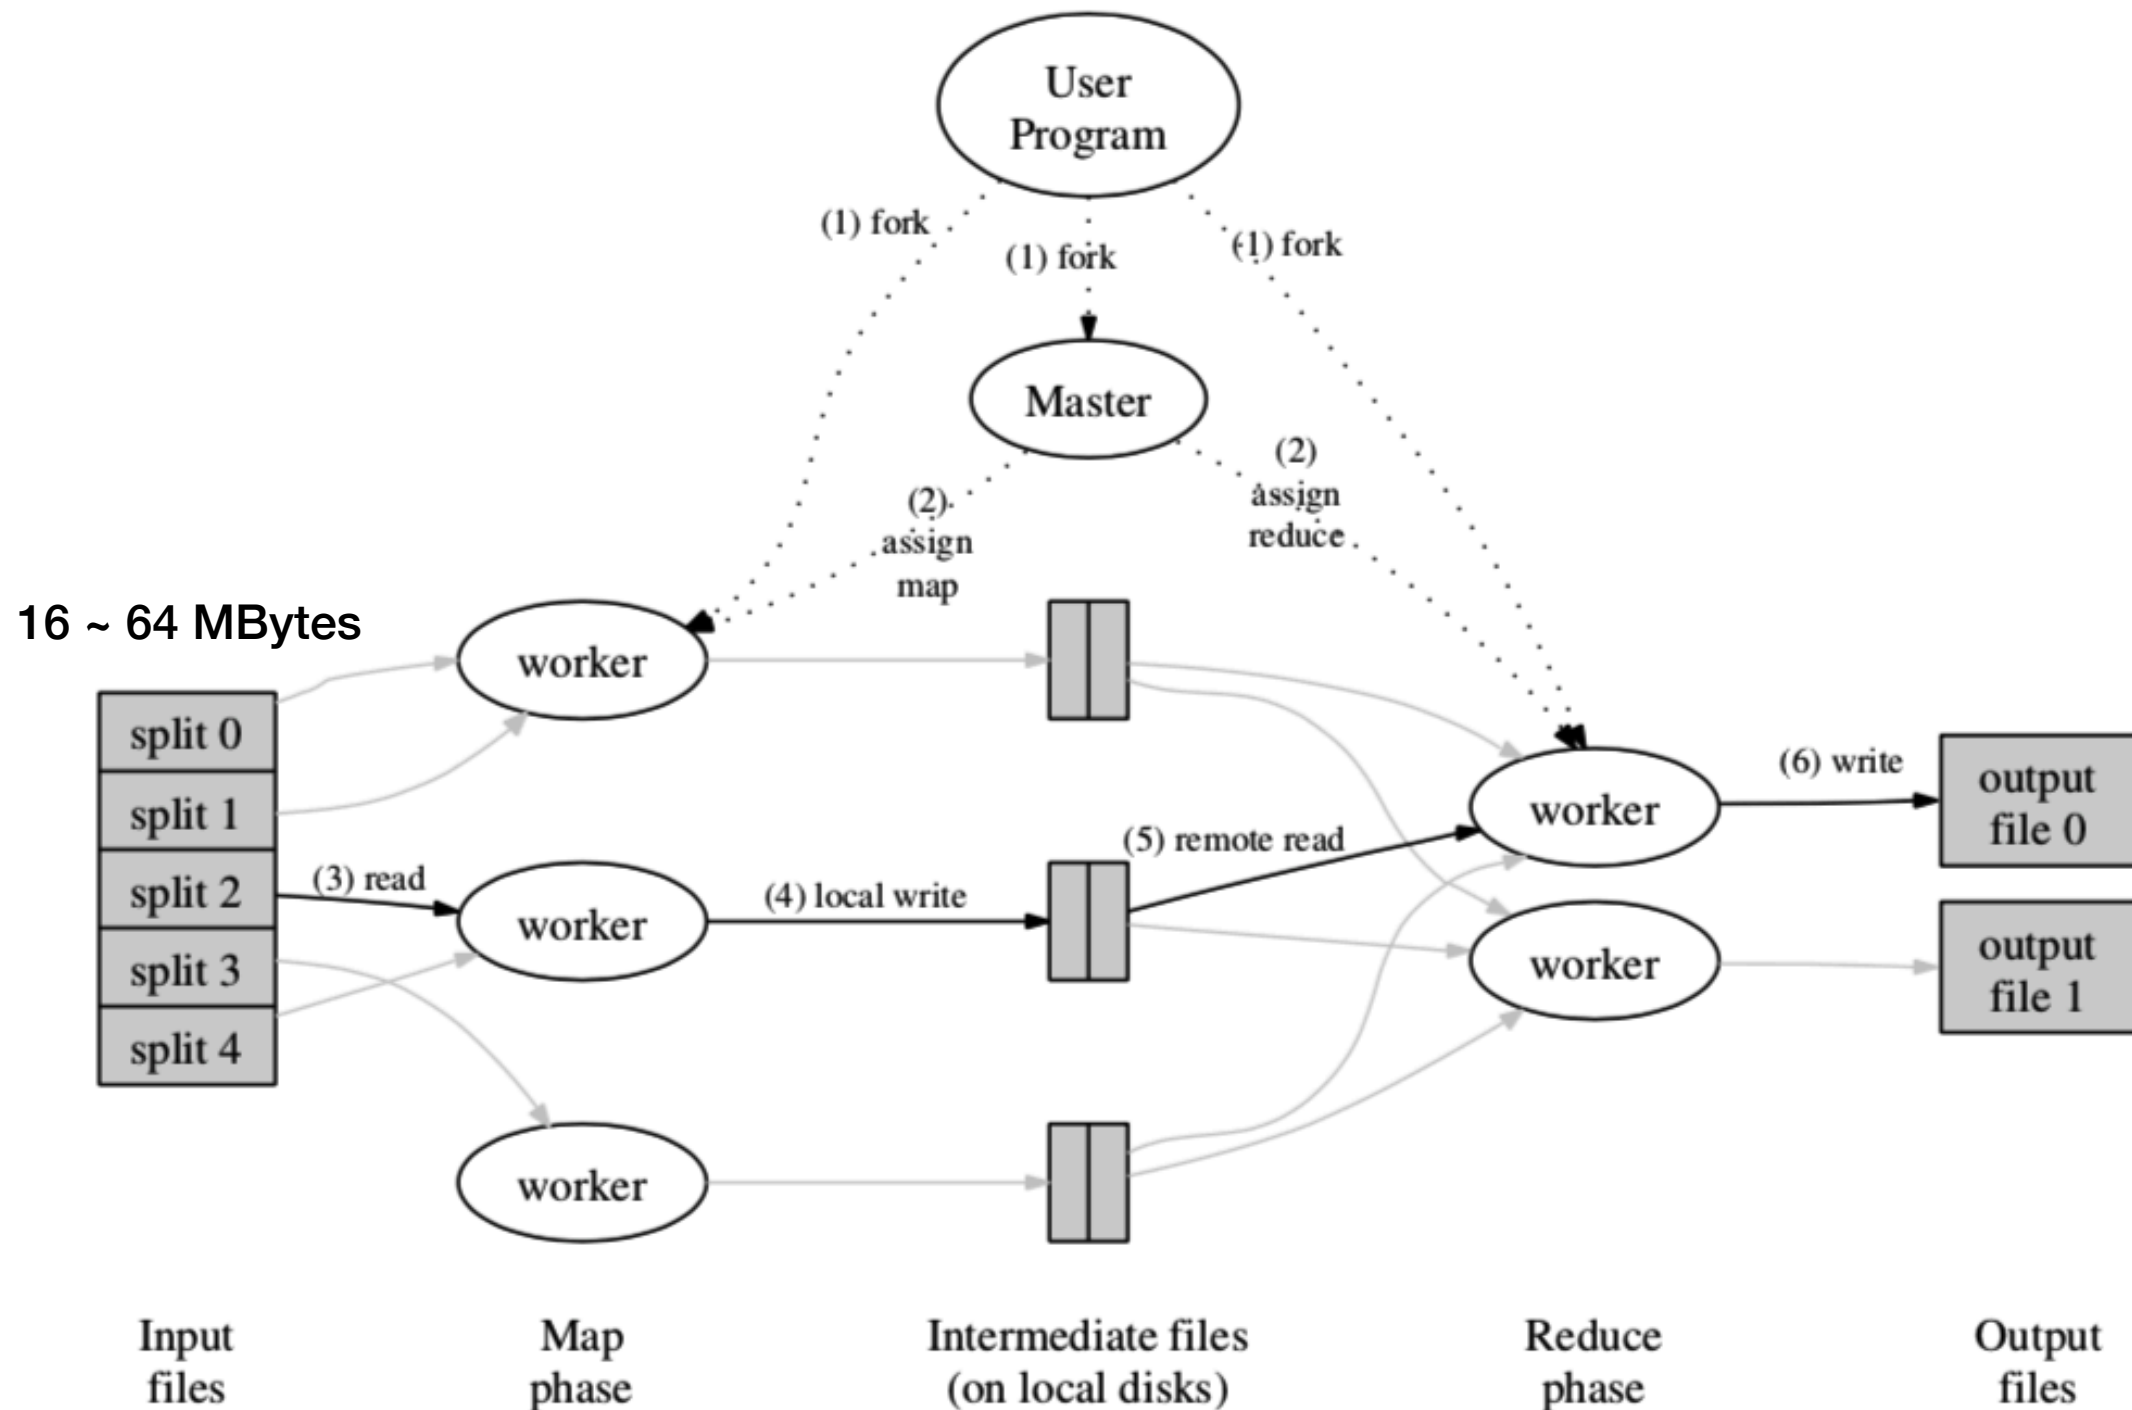
**OS: Linux**

**RAM: 2-4 GB**

**Storage: inexpensive IDE disks**



# Execution overview



# Worker failure tolerance

- \* Every failed worker and map task is reset to Idle state by Master for rescheduling.**
- \* Completed map tasks are re-executed on a failure while completed reduce tasks aren't.**
  - writing to local disks or global file system.**

# Master failure tolerance

- \* Writing periodic checkpoints of the master data structures.**
- \* New copy from the last checkpoint state of Master**
  - using this, master can be restarted if master dies.**

# Locality

- \* To minimize network bandwidth, Replication of the input data block is used.**
  - master schedules a map task on a machine that contains a replica or is near a replica.**

# Backup tasks

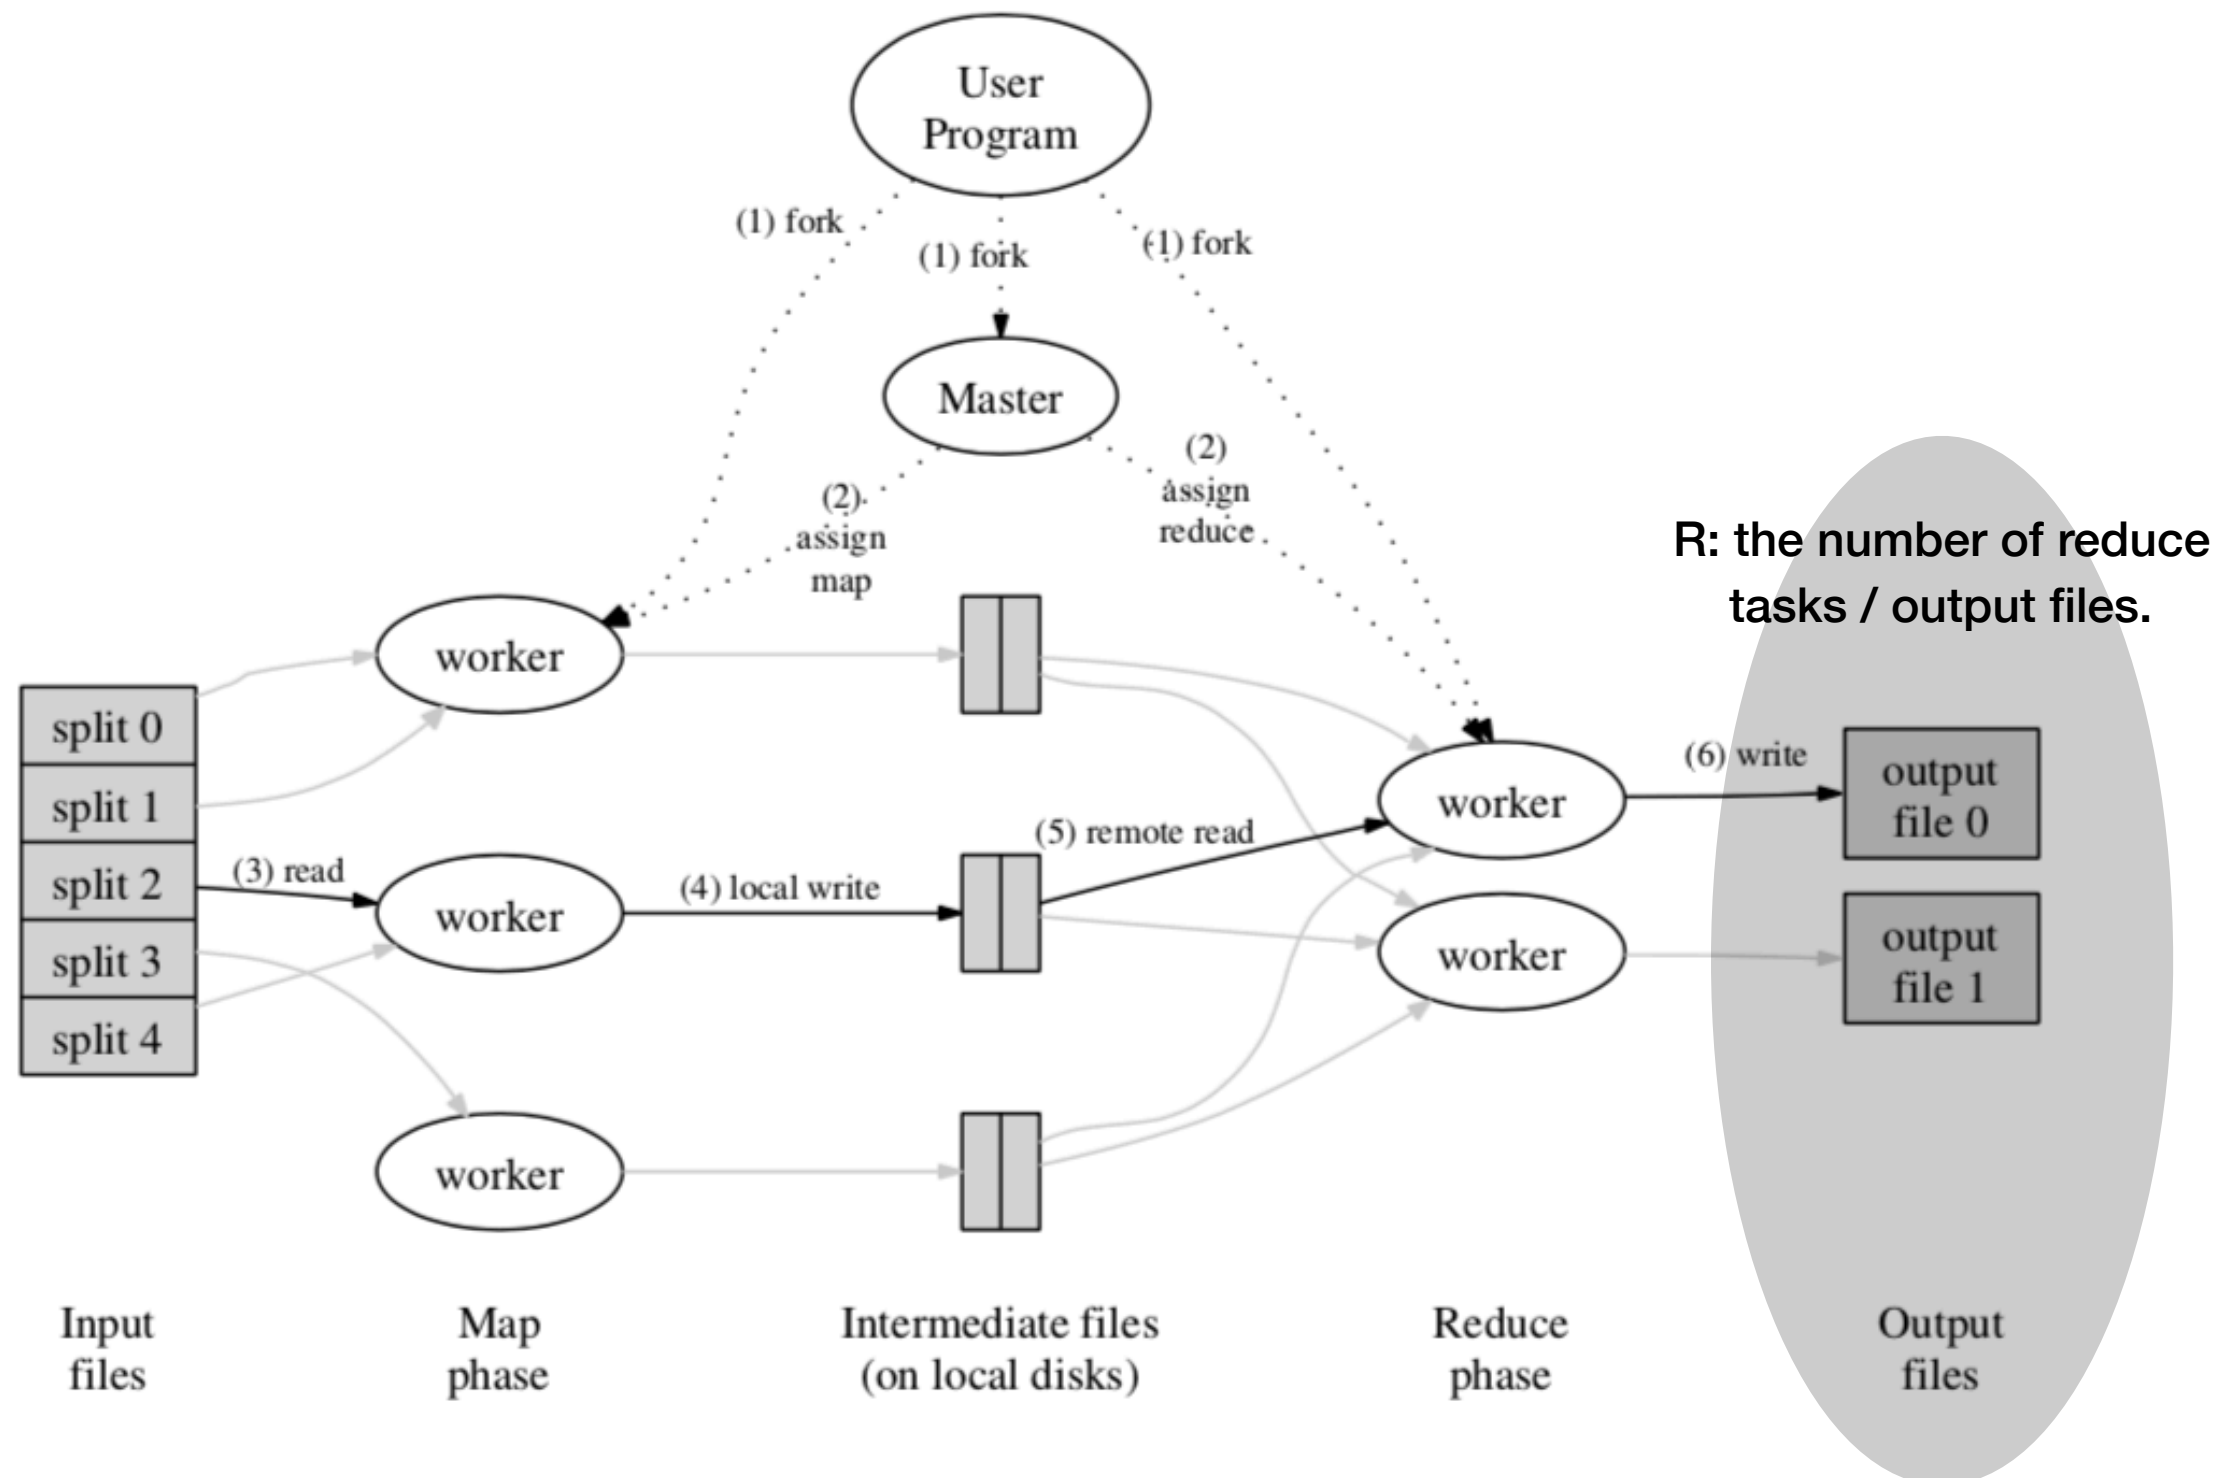
- \* When MapReduce operation is close to completion, the master schedules backup executions of the remaining in-progress tasks.
  - Because remaining in progress tasks can be **stragglers**, which can increase overall run time.

sorting program takes 44% shorter to complete using Backup tasks!

# Refinements



# Partitioning Function



# Partitioning Function

- \* Default partitioning function :  $\text{Hash}(\text{key}) \bmod R$
- \* special partitioning function:  $\text{Hash}(\text{Hostname}(\text{urlKey}) \bmod R$ 
  - In specific cases (e.g. all entries for a single host to end up in the same output file)



# Combiner Function

- \* Partial merging of thousands of records of intermediate key value pair
  - Written to an intermediate file before reduce task.
- \* Combiner function code and Reduce function code is exactly the same.

**Significantly speeds up!**

# Counters

```
Counter* uppercase;  
uppercase = GetCounter("uppercase");  
  
map(String name, String contents):  
    for each word w in contents:  
        if (IsCapitalized(w)):  
            uppercase->Increment();  
            EmitIntermediate(w, "1");
```

- A facility to count occurrences of various events.
- Master aggregates the counter values from successful map and reduce tasks.

**Useful for sanity checking the behavior of MapReduce operations.**

# Results

# Cluster configuration

## **1800 commodity pcs**

- **2GHz Intel Xeon cpu with hyper threading**
- **4GB RAM**
- **160GB IDE disks**
- **Gigabit Ethernet link**

# Performance

input data split: 64MB \* 15000  
output file: 1




Figure 2: Data transfer rate over time

**Performance Measure on Grep program**  
(Searching for a particular pattern in one terabyte of data).

# Performance

input data split: 64MB \* 15000  
output file: 4000

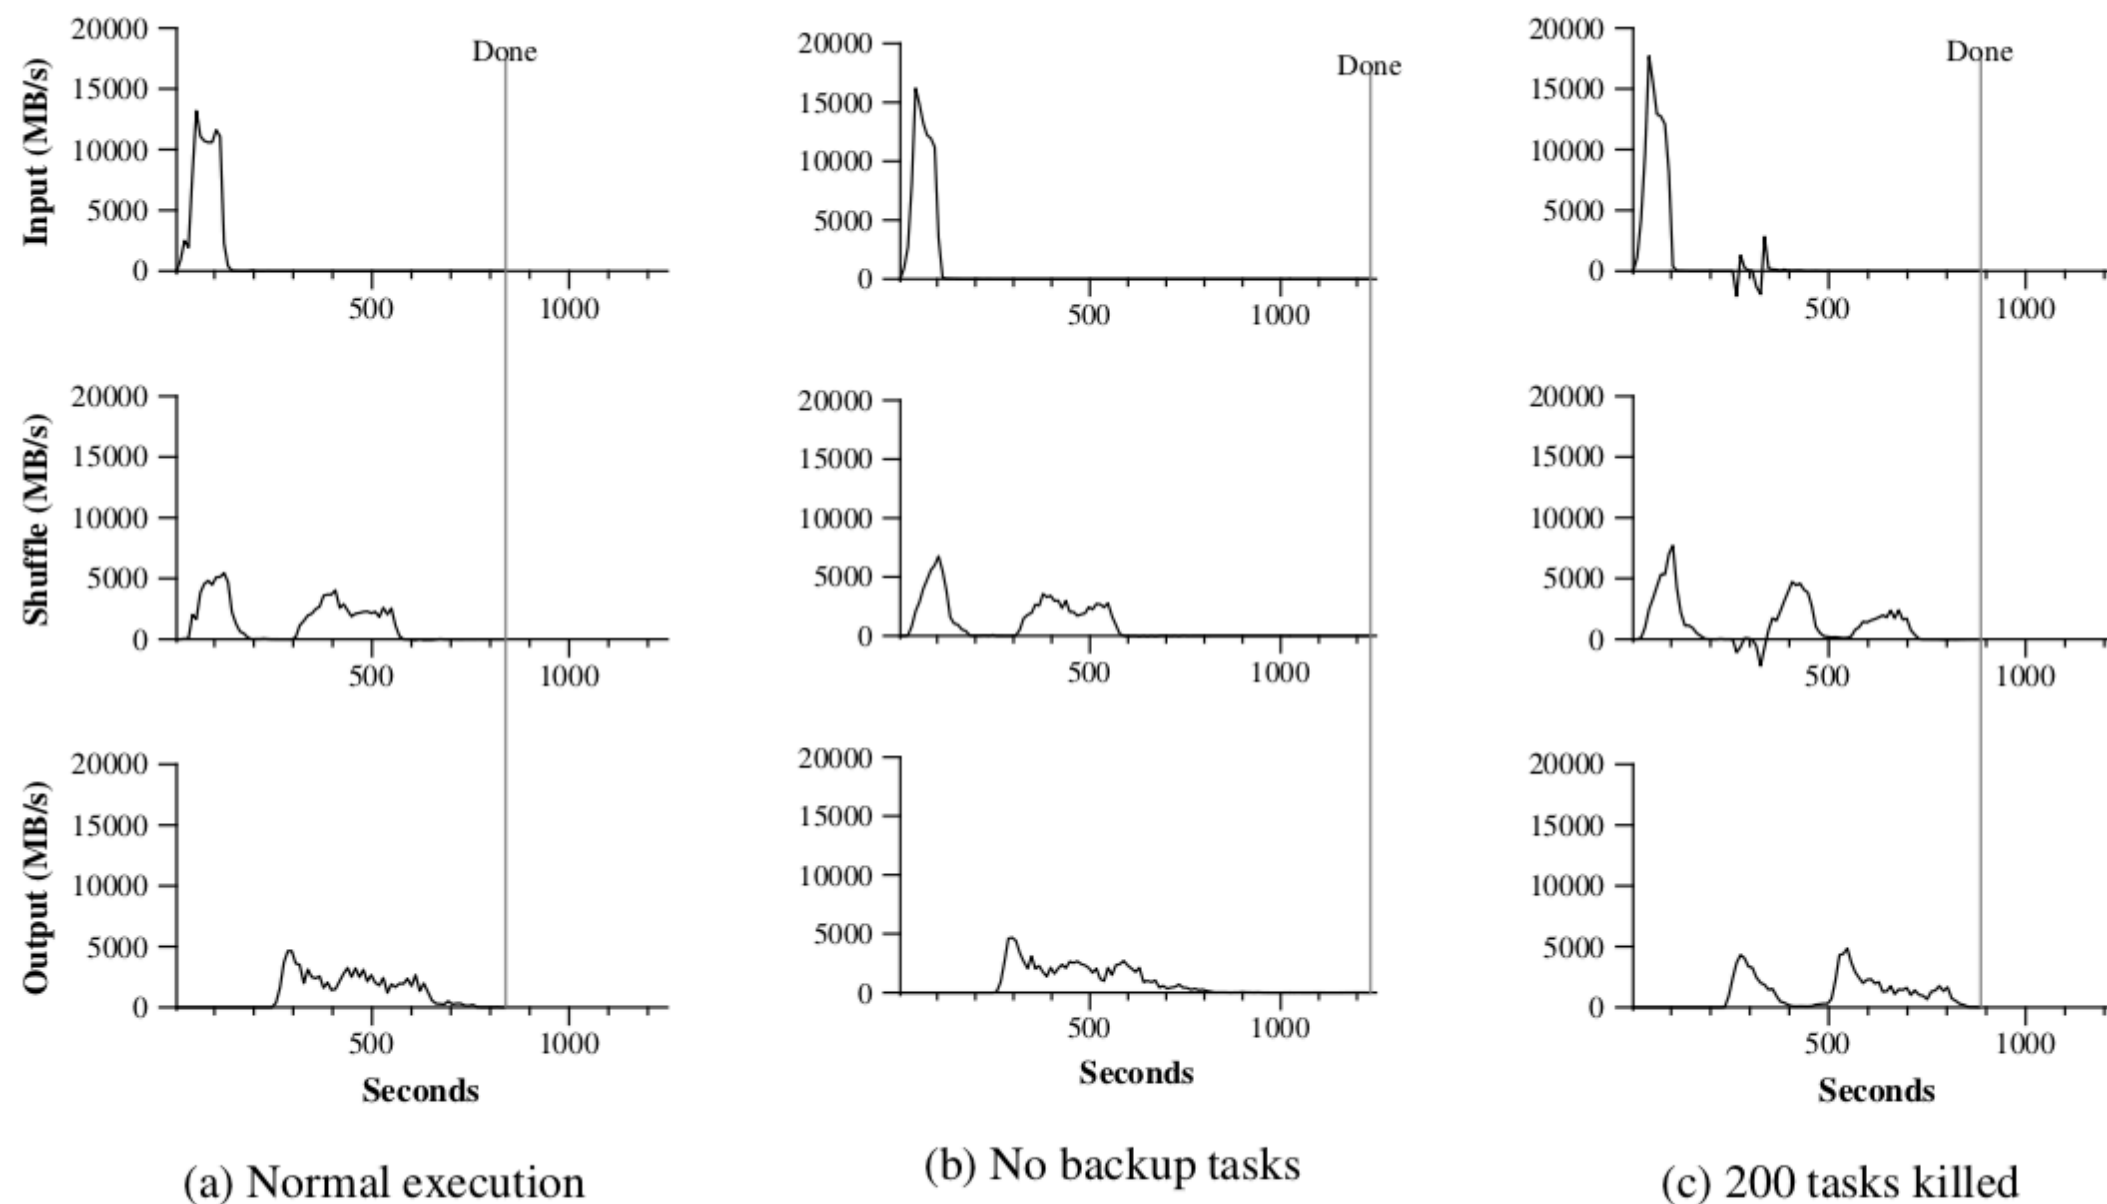


Figure 3: Data transfer rates over time for different executions of the sort program

**Performance Measure on Sorting one terabyte of data.**

# Conclusion



# Conclusion

- \* An programming model easy to use
  - hiding the details of parallelization, fault-tolerance, locality optimization.
- \* Large variety of problems are easily expressible.
  - Google web search service, data mining...
- \* Computation scaling on large clusters.
  - efficient to use machine resources.



**THNAKS**